

SUMMARY

Data Engineer with a software engineering background and 3+ years focused on data engineering and analytics platforms. Experienced building production systems with **dbt**, **Snowflake**, **BigQuery**, **Apache Airflow**, Python, and SQL across private equity, SaaS, retail, and fintech. Focused on end-to-end platform ownership, warehouse performance, data quality, orchestration, and analytics engineering.

EXPERIENCE

- Data Engineer – ECCO

07/2026–Present

 - Leading end-to-end discovery and design of a centralized HR data platform spanning European business units, translating requirements from HR and IT stakeholders across multiple countries into data ownership, integration, and reporting requirements
 - Designing Airflow-orchestrated ingestion and **dbt** silver/gold layers for workforce data across European operations, providing standardized datasets for analytics and Power BI reporting
 - Prototyping a governed AI and semantic layer with **Snowflake Cortex**, Semantic Views, Cortex Agents, and Cortex Search for natural-language analytics and RAG over internal documents stored in object storage
- Data Engineer – Lortech Solutions

08/2024–07/2026

 - Delivered data and analytics platforms across **7 client engagements** in private equity, SaaS, retail, and fintech, processing multi-terabyte daily workloads across Snowflake and BigQuery
 - Built and owned a business-critical investor relations platform with **350+ dbt models** and **350+ automated tests**, integrating Salesforce, Outlook, Box, FTP, REST APIs, and SOAP services through Airflow
 - Reduced a high-cost Snowflake workload from **55 minutes to 30 seconds** and lowered BigQuery spend by **35%** through SQL redesign, warehouse tuning, partitioning, materialization changes, and model refactoring
 - Standardized source onboarding through metadata-driven dbt model generation, reusable macros, YAML templates, automated validation, and CI/CD checks
 - Improved reliability of multi-tenant Airflow environments through freshness and source-availability checks, run-duration monitoring, SLA validation, and automated Slack/email incident reporting
 - Standardized deployments for Kubernetes-based Airflow environments using versioned Docker images, AWS ECR, and CI pipelines
- Independent Software & Data Consultant – Self-employed

11/2018–04/2025

 - Delivered backend systems, scraping workflows, data extraction pipelines, and automation tools for clients across European and US markets
 - Built ingestion and processing systems with Python, PostgreSQL, Redis, browser automation, and reverse-engineered APIs, including deduplication, retry handling, and schema normalization for unstable external sources
- Python Developer – Decision Sciences Company

11/2022–05/2023

 - Developed Python data workflows, API integrations, and automation for a production AI pricing platform in the insurance domain

TECHNICAL PROJECTS

- dbt-doctor – GitHub

Open-source CLI for dbt project validation, detecting broken refs, missing sources, untested models, and unsafe incremental configurations. Runs locally and in GitHub Actions; **65+ stars**.
- moneySense-data – GitHub

Event-driven pipeline for Polymarket market data using Kafka and ClickHouse for the hot path and Parquet on MinIO for the cold archive.

SKILLS

Data Engineering: **dbt**, **Apache Airflow**, **Snowflake**, **BigQuery**, **Databricks**, **Apache Spark**, ELT/ETL, dimensional modeling, data quality, observability, warehouse performance and cost optimization, Power BI

Programming: **Python**, **SQL**, Go, TypeScript

Infrastructure: **AWS**, **GCP**, **Docker**, **Kubernetes**, Kafka, PostgreSQL, ClickHouse, MinIO, Parquet, Iceberg, Trino, CI/CD

EDUCATION
Polish-Japanese Academy of Information Technology (PJATK)
B.Sc. Computer Science
2024–2028

CERTIFICATIONS
SnowPro Core (Snowflake)
Credential ID: 156908869
dbt Developer (dbt Labs)
Credential ID: 135541190

PUBLICATIONS
Wide Tables in Snowflake
AI Analytics Workflows
Salesforce dbt Automation
Warehouse Cost Attribution