

Joachim Hodana

SENIOR DATA ENGINEER

✉ me@joachimhodana.com 🌐 joachimhodana.com 🌐 joachim-hodana 🌐 joachimhodana
📍 Warsaw, Poland ☎ +48 573 918 052

SUMMARY

Data engineer with eight years in production systems, the last two spent full-time on the modern data stack. Mostly dbt and Airflow on Snowflake and BigQuery, building for enterprise clients in private equity, retail, fintech and SaaS - raw ingestion through to the models their BI tools sit on. Happy owning a platform end to end, including the parts nobody volunteers for: test coverage, warehouse cost, and documentation the next engineer will actually open.

EXPERIENCE

Senior Data Engineer – ECCO (contract) 07/2026–Present

- Data engineering for the group’s retail and supply chain analytics, building dbt models and the Power BI datasets the business reports on

Data Engineer – Lortech Solutions 08/2024–07/2026

- Delivered data platforms for seven enterprise clients across private equity, retail, fintech and SaaS, from raw ingestion up to the semantic layer their BI tools read
- Owned an investor relations platform end to end: around 350 dbt models with matching test coverage, fed by Airflow pipelines pulling from Salesforce, Outlook, Box, FTP and a pair of legacy SOAP services
- Cut one Snowflake workload from 55 minutes to 30 seconds by rewriting the query and resizing the warehouse behind it, and took 35% off a client’s BigQuery bill through partitioning and materialization changes
- Wrote a metadata-driven generator for dbt models, so onboarding a new Salesforce object became a YAML edit instead of a few hundred hand-written lines of SQL
- Added freshness, source availability and run-duration monitoring to a multi-tenant Airflow estate, reporting into Slack so failures surfaced before the client noticed them
- Moved Airflow onto Kubernetes with versioned Docker images in ECR and CI to promote them, replacing instances that had been configured by hand

Independent Software & Data Consultant – Self-employed 11/2018–07/2026

- Scoped, priced and delivered projects for clients in Poland, the EU, US and UK, usually as the only engineer on the work
- Backend systems, APIs and automation in Python, Go and TypeScript, on PostgreSQL and Redis; more recently RAG pipelines and LLM tooling for clients automating research and support
- Data extraction for clients who needed data no vendor exposed: scraping, browser automation and reverse-engineered APIs, which meant handling deduplication, retries and schema drift as a matter of course

Python Developer – Decision Sciences Company 11/2022–05/2023

Python data workflows, API integrations and automation for an insurance AI product.

TECHNICAL PROJECTS

dbt-doctor – GitHub

Open-source CLI that scores a dbt project and flags what passes `dbt test` but breaks in production: broken refs, missing sources, incremental models without a unique key. Runs as a GitHub Action. Around 700 downloads.

moneysense-data – GitHub

Event-driven pipeline for Polymarket market data. Kafka into ClickHouse for the hot path, Parquet on MinIO for the cold archive.

Dealscan

Aggregates car listings from a dozen sources that each define a listing differently. Browser automation and reverse-engineered APIs in, deduplicated PostgreSQL out.

SKILLS

Data Engineering: dbt, Apache Airflow, Snowflake, BigQuery, ELT/ETL, dimensional modeling, data quality & testing, observability, warehouse cost optimization, Power BI

Programming: Python, SQL, Go, TypeScript

Infrastructure: AWS, GCP, Docker, Kubernetes, CI/CD, PostgreSQL, ClickHouse, Kafka, MinIO, Parquet

EDUCATION

Polish-Japanese Academy of Information Technology (PJATK)
B.Sc. Computer Science
2024–2028

CERTIFICATIONS

SnowPro Core (Snowflake)
Credential ID: 156908869
dbt Developer (dbt Labs)
Credential ID: 135541190

PUBLICATIONS

Wide Tables in Snowflake
AI Analytics Workflows
Salesforce dbt Automation
Warehouse Cost Attribution