

SUMMARY

Senior Data Engineer with eight years of experience building production software and data systems, including more than two years focused on enterprise analytics platforms. Delivered end-to-end ELT pipelines, warehouse models, and orchestration workflows across **Snowflake** and **BigQuery** for clients in private equity, SaaS, retail, and fintech. Hands-on with **dbt**, **Apache Airflow**, Python, SQL, Docker, Kubernetes, AWS, and GCP, with a focus on performance and cost optimization, data quality, observability, API ingestion, and maintainable delivery.

EXPERIENCE

- Senior Data Engineer – ECCO07/2026–Present
- Building a data platform for HR analytics and reporting, working directly with IT and HR stakeholders to translate business requirements into reliable datasets and reporting foundations
- Data Engineer – Lortech Solutions08/2024–07/2026
- Delivered analytics engineering systems across **7+ client projects** in private equity, SaaS, retail, and fintech, covering **1000+ dbt models**, **100+ orchestration DAGs**, and multi-terabyte daily workloads across Snowflake and BigQuery
 - Owned end-to-end development of a business-critical investor relations platform with **350+ dbt models**, **350+ tests**, and Airflow pipelines integrating Salesforce, Outlook, Box, FTP, REST APIs, and SOAP services
 - Reduced one high-cost Snowflake workload from **55 minutes to 30 seconds** and lowered BigQuery spend by **35%** through SQL redesign, warehouse tuning, partitioning, materialization changes, and model refactoring
 - Built metadata-driven dbt model generation and validation workflows using reusable macros, YAML templates, model tests, and CI/CD checks, reducing repetitive SQL and standardizing new source onboarding
 - Added observability and SLA validation to multi-tenant Airflow environments, including freshness checks, source availability monitoring, run-duration tracking, and Slack/email incident reporting
 - Standardized deployment workflows for Kubernetes-based Airflow environments using versioned Docker images, AWS ECR, and CI pipelines
- Independent Software & Data Consultant – Self-employed11/2018–04/2025
- Delivered backend systems, scraping workflows, data extraction pipelines, and automation tools for clients across EU and US markets
 - Built ingestion and processing systems using Python, PostgreSQL, Redis, browser automation, and reverse-engineered APIs, including deduplication, retry handling, and schema normalization for unstable external sources
- Python Developer – Decision Sciences Company11/2022–05/2023
- Developed Python data workflows, API integrations, and automation for a production AI pricing product in the insurance domain

TECHNICAL PROJECTS

- dbt-doctor – GitHub
- Open-source CLI for dbt project validation, detecting broken refs, missing sources, untested models, and unsafe incremental configurations. Runs locally and in GitHub Actions; **700+ downloads**.
- moneysense-data – GitHub
- Event-driven pipeline for Polymarket market data, using Kafka and ClickHouse for the hot path and Parquet on MinIO for the cold archive.
- Dealscan
- Multi-source automotive ingestion system using browser automation and reverse-engineered APIs to normalize and deduplicate heterogeneous listings in PostgreSQL.

SKILLS

Data Engineering: dbt, Apache Airflow, Snowflake, BigQuery, ELT/ETL, dimensional modeling, data quality and testing, observability, warehouse cost optimization, Power BI

Programming: Python, SQL, Go, TypeScript

Infrastructure: AWS, GCP, Docker, Kubernetes, CI/CD, PostgreSQL, ClickHouse, Kafka, MinIO, Parquet

EDUCATION

Polish-Japanese Academy of Information Technology (PJATK)

B.Sc. Computer Science

2024–2028

CERTIFICATIONS

SnowPro Core (Snowflake)

Credential ID: 156908869

dbt Developer (dbt Labs)

Credential ID: 135541190

PUBLICATIONS

Wide Tables in Snowflake

AI Analytics Workflows

Salesforce dbt Automation

Warehouse Cost Attribution